

Cost Function Design & Markov Chains with PERT

Agentic Congress Planner (ACP) - Architecture Document
Prepared by Claudia | Sunrise Gen AI | Feb 2026

Cost Function Design & Markov Chains with PERT

Part A: Cost Function Design

1. Multi-Objective Cost Function

The planner optimizes across competing objectives. The composite cost function:

$C_{total} = w1*C_{schedule} + w2*C_{resource} + w3*C_{risk} + w4*C_{quality} + w5*C_{disruption}$

Where weights $w1..w5$ are configurable per congress priority (tuned via historical outcomes).

2. Individual Cost Components

C_schedule -- Schedule Cost

```
C_schedule = SUM over all tasks [  
    alpha * max(0, actual_end - due_date)^2      // Tardiness penalty (quadratic)  
    + beta * max(0, due_date - actual_end)        // Earliness bonus (linear, diminishing)  
    + gamma * (is_critical_path * tardiness_penalty) // Critical path multiplier  
]
```

- Quadratic tardiness: Small delays are tolerable, large delays are catastrophic (non-linear penalty)
- Critical path multiplier (gamma): Delays on CP tasks penalized 3-5x more
- MS Planner mapping: dueDateTime vs simulated completedDateTime

C_resource -- Resource Utilization Cost

```
C_resource = SUM over all assignees [  
    delta * max(0, utilization - U_max)^2      // Over-allocation penalty  
    + epsilon * max(0, U_min - utilization)     // Under-utilization waste  
    + zeta * context_switch_count                // Multi-tasking penalty  
]
```

- Derived from: assignments field -- count concurrent tasks per person per time window
- Historical calibration: Compute each person's U_{max} from past congresses
- Context switch cost: Each additional parallel task adds ~20% overhead

C_risk -- Risk Exposure Cost

```
C_risk = SUM over all tasks [  
    eta * P(delay > threshold) * impact_magnitude  
]  
- P(delay): From Monte Carlo simulation output
```

- Impact magnitude: Derived from priority field + downstream dependency count in DAG
- High fan-out tasks (many dependents) get higher impact scores

C_quality -- KOL & Speaker Alignment Quality

```
C_quality = SUM over all speaker assignments [
    theta * topic_mismatch_penalty
    + iota * preferenceViolation_penalty
    + kappa * audience_overlap_penalty
]
- Topic mismatch: NLP similarity between speaker expertise and assigned topic
- Preference violations: Soft constraint from taskDetail.description
- Audience overlap: If two sessions share >40% target audience, penalize parallel scheduling
```

C_disruption -- Reactive Replanning Cost

```
C_disruption = SUM over all replan events [
    lambda * cascade_depth
    + mu * human_approval_latency
    + nu * plan_delta
]
- Stability premium: Users prefer plans that don't change drastically
- Cascade depth: From dependency DAG traversal
```

3. Cost Function Calibration from Historical Data

1. Compute actual costs from historical congresses (delays, overruns, speaker issues)
2. Regression: fit weights w1..w5 to minimize |predicted_cost - actual_outcome_cost|
3. Output: Validated weight vector for current year planning

Part B: Markov Chains + PERT Integration

1. Why PERT Alone Is Insufficient

Standard PERT assumes:

- Beta distribution for task durations (optimistic, most likely, pessimistic)
- Independence between tasks (only structural dependency via DAG)
- Static estimates -- no state transitions during execution

Problems for congress planning:

- Task states evolve -- a task "in progress" for too long signals trouble
- External events change the probability landscape
- Resource states matter -- overloaded person affects ALL their tasks

2. Markov Chain Enhancement: State-Based Task Modeling

Each task is modeled as a Markov chain with states:

States S = {Not Started, Planning, In Progress, Blocked, Under Review, Completed, Cancelled}

Mapping to MS Planner:

- Not Started: percentComplete = 0, no assignedDateTime

- Planning: percentComplete = 0, has assignedDateTime
- In Progress: percentComplete = 50
- Blocked: percentComplete = 50 + no progress for > T_block threshold
- Under Review: percentComplete = 50 + checklist mostly done
- Completed: percentComplete = 100
- Cancelled: task deleted or marked in description

3. Transition Probability Matrix (from Historical Data)

Each cell $P[i,j]$ represents probability of transitioning from state i to state j per time period.
Calibrated per task category, assignee historical performance, and congress phase.

Example transitions:

- Not Started -> Planning: 0.7
- Planning -> In Progress: 0.8
- In Progress -> Under Review: 0.4
- In Progress -> Blocked: 0.15
- Under Review -> Completed: 0.7
- Blocked -> In Progress: 0.6
- Completed and Cancelled are absorbing states

4. PERT-Markov Hybrid Model

Standard PERT:

Duration $\sim \text{Beta}(a, m, b)$ where a=optimistic, m=most_likely, b=pessimistic

Enhanced PERT-Markov:

At each time step t:

1. Current state s_t of each task (from Markov chain)
2. Time spent in current state: τ
3. Conditional duration estimate:
 $D_{\text{remaining}} | s_t, \tau \sim \text{Updated_Beta}(a', m', b')$
 where a', m', b' are adjusted based on:
 - Current state (blocked tasks get inflated pessimistic)
 - Dwell time (longer in IP -> right-shift distribution)
 - Resource state (assignee overloaded -> inflate)

5. Absorbing Markov Chain Analysis

Fundamental Matrix: $N = (I - Q)^{-1}$

Where Q = transient state transition submatrix

$N[i,j] =$ expected number of times process visits state j, starting from state i

Expected time to completion from state i = $\text{SUM}(N[i,:])$

This gives us:

- Expected task duration accounting for blocked/review loops
- Probability of cancellation vs completion from any state
- Variance of completion time (from N matrix properties)

6. Markov-Modulated PERT Network

Traditional PERT Network:

Task A (duration_A) -> Task B (duration_B) -> Task C (duration_C)

Markov-Modulated:

Task A [MC_A state] -> Task B [MC_B state] -> Task C [MC_C state]

Each task's duration is conditioned on:

1. Its own Markov chain state
2. Predecessor's terminal state (completed normally vs rushed vs partial)
3. Shared resource Markov chains (person availability state)

Dependency types become richer:

- Finish-to-Start (FS): Standard -- B starts when A completes
- State-conditional FS: B's duration changes based on HOW A completed
- Resource-coupled: A and B share assignee -> joint Markov chain

7. Combined Simulation Algorithm

Initialize:

For each task: set state = Not Started

For each task: fit PERT params (a,m,b) from historical data

Build transition matrices per task category

For each simulation run (1..10000):

t = 0

While not all tasks absorbed:

t += 1

For each active task:

```
new_state = sample(transition_matrix[current_state])
```

If entering "In Progress" and not duration_sampled:

```
duration = sample(Beta(a, m, b))
```

If assignee.state == Overloaded:

```
duration *= resource_penalty_factor
```

If new_state == Blocked:

```
block_time = sample(Exponential(lambda_block))
```

```
duration += block_time
```

Check completion: if time_in_progress >= duration -> transition

Inject external disruptions at time t

Record: total_time, critical_path, cost_function_value

Output: distributions, percentiles, risk metrics

8. When to Use What

- Pure PERT: Quick estimates, well-understood tasks (no state awareness)
- Pure Monte Carlo: Duration uncertainty exploration (no state transitions)
- Markov Chain: State-based progression modeling (memoryless assumption)
- PERT + Markov + MC: Full system -- state-aware, uncertainty-modeled, dependency-coupled

9. Addressing the Memoryless Limitation

Standard Markov chains are memoryless. For congress planning, dwell time matters. Solutions:

- Semi-Markov Process: Transition probabilities depend on time spent in current state
- Phase-type distributions: Represent non-exponential durations as sequences of exponential phases
 - Practically: Discretize time into periods (e.g., daily). Transition matrix changes based on days_in_state -- this is a time-inhomogeneous Markov chain